

CUDA-OpenGL interoperability

Visualization

Sergiu Oprea
Albert García García

soprea @ dtic.ua.es
agarcia @ dtic.ua.es

Contents

Introduction

OpenGL

CUDA + OpenGL interop

Example

Introduction

Computer graphics applications

Visualization

- Graphics-oriented hardware
- Use CUDA to perform calculations and then display visual results
- Efficient compute-visualization pipeline
 - No extra memory movements involved
 - Unlocks CPU cycles for other application tasks

OpenGL

OpenGL

Open Graphics Library

- OpenGL is an interface for the graphic hardware programmer
- API:
 - Primitives: points, lines and polygons
 - Properties: colors, lighting, textures
 - View: camera position and perspective

OpenGL

Example

```
glBindTexture( GL_TEXTURE_2D, textureID);  
glColor3f(1.0f,0,0);  
glBegin(GL_QUADS);  
    glTexCoord2i( 0, h);  
    glVertex3f(0,0,0);  
  
    glTexCoord2i(0,0);  
    glVertex3f(0,1.0f,0);  
  
    glTexCoord2i(w,0);  
    glVertex3f(1.0f,1.0f,0);  
  
    glTexCoord2i(w,h);  
    glVertex3f(1.0f,0,0);  
glEnd();  
  
SwapBuffers(hDC);
```

OpenGL is *state-based*,
parameters are sticky.

CUDA + OpenGL interop

CUDA + OpenGL

Interoperability

- CUDA for calculations, data generation, image manipulation,...
- OpenGL to draw the pixels & vertices on the screen
- Efficient interaction/coordination
 - Share data via a common memory in the framebuffer

CUDA + OpenGL

Interoperability

- CUDA C uses memory management techniques similar to those of C (malloc, pointers,...)
- OpenGL stores data in generic and abstract buffers called buffer objects.
- The interaction between CUDA/OpenGL is based on a simple concept:
 - Map/Unmap an OpenGL buffer for CUDA interoperability

Example

CUDA + OpenGL

Example

Steps to configure OpenGL with CUDA:

1. Create a GL window
2. Create a GL context
3. Configure the GL viewport and coordinate system
4. Create a CUDA context
5. Generate one or more GL buffers to share with CUDA
6. Register buffers on CUDA

CUDA + OpenGL

Example

Steps to configure OpenGL with CUDA:

1. Create a window (OpenGL)

- Every S.O. does it differently, in WIN 32 it is:
 - `CreateWindowEx()` Win32
 - It needs a Windows HDC:
`HDC hDC;`
`hDC=GetDC(hWnd)`
- We can use GLUT/GLFW to manage the creation/destruction of windows.

CUDA + OpenGL

Example

Steps to configure OpenGL with CUDA:

2. Create a GL context

```
// create a wGL rendering context
HGLRC hGLRC;
hGLRC=wglCreateContext(hDC);
// Enable rendering context
wglMakeCurrent(hDC,hGLRC);
// load OpenGL extensions to use buffers
glewInit();
```

CUDA + OpenGL

Example

Steps to configure OpenGL with CUDA:

3. Configure the GL viewport and the coordinate system

```
// configure viewport size
glViewport(0, 0, width, height);
// setup projection mode
glMatrixMode(GL_PROJECTION);
glLoadIdentity();
glOrtho(0,1.0f,0,1.0f,-1.0f,1.0f);
glMatrixMode(GL_MODELVIEW);
glLoadIdentity();
```

CUDA + OpenGL

Example

Steps to configure OpenGL with CUDA:

3. Configure the GL viewport and the coordinate system

.....

```
// Enable depth test
glEnable(GL_DEPTH_TEST);
// Clean color and viewport
glClearColor(1.0f, 1.0f, 1.0f, 1.5f);
glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT);
```


CUDA + OpenGL

Example

Steps to configure OpenGL with CUDA:

4. Create a CUDA context

- OpenGL context must be created first
- Create a CUDA context:
 - Driver API: Use `cuGLCtxCreate()` instead of `cuCtxCreate()`
 - Runtime API: use `cudaGLSetGLDevice()` before using the GL API
- The CUDA/OpenGL interaction is defined in :
 - `cudaogl.h` for the Driver API
 - `cuda_gl_interop.h` CUDA C Runtime

CUDA + OpenGL

Example

Steps to configure OpenGL with CUDA:

5. Create OpenGL buffers

```
GLuint bufferID;  
// buffer ID  
glGenBuffers(1,&bufferID);  
// set UNPACK for the current buffer  
glBindBuffer(GL_PIXEL_UNPACK_BUFFER, bufferID);  
// memory assignment  
glBufferData(GL_PIXEL_UNPACK_BUFFER, width * height * 4,  
NULL, GL_DYNAMIC_COPY);
```

CUDA + OpenGL

Example

Steps to configure OpenGL with CUDA:

6. Registrar los buffers en CUDA

- Driver API:

- `cuGLRegisterBufferObject(GLuint bufferobj);`
- `cuGLUnregisterBufferObject(GLuint bufferobj);`

- Runtime API:

- `cudaGLRegisterBufferObject(GLuint bufObj);`
- `cudaGLUnregisterBufferObject(GLuint bufObj);`

CUDA + OpenGL

Example

Map/Unmap

```
// map
float4 *dptr;
cudaGraphicsMapResources(1, vbo_resource, 0);
size_t num_bytes;
cudaGraphicsResourceGetMappedPointer((void **)&dptr, &num_bytes,
                                     *vbo_resource));

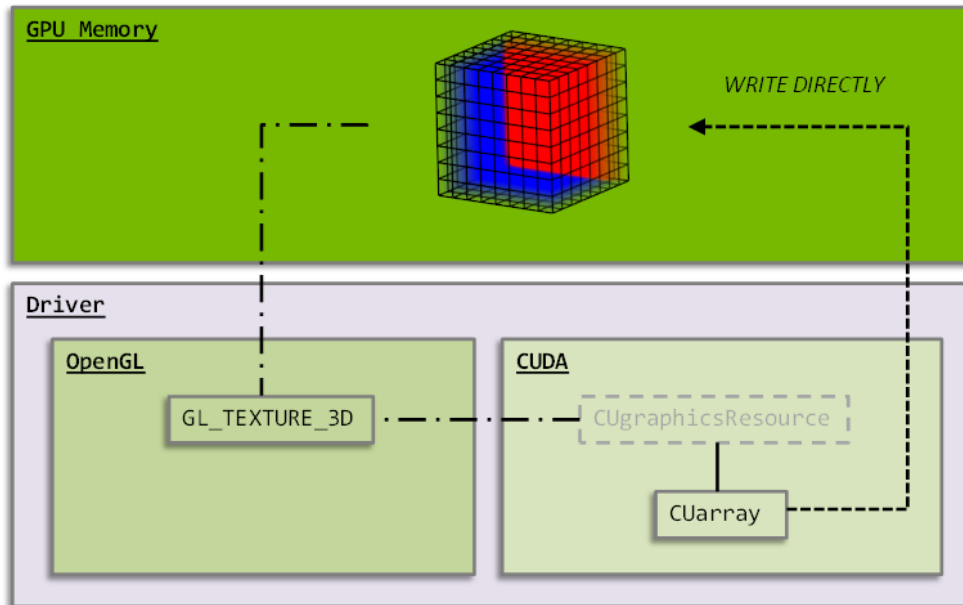
// cuda kernel
launch_kernel(dptr, mesh_width, mesh_height, g_fAnim);

// unmap
cudaGraphicsUnmapResources(1, vbo_resource, 0);

// draw
Draw()
```

CUDA + OpenGL

Example Map/Unmap



https://github.com/nvpro-samples/gl_cuda_interop_pingpong_st

CUDA-OpenGL interoperability

Visualization

Thanks for your attention!

These slides have been modified/remixed using the TeachingKit licensed by NVIDIA and the University of Illinois under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).



Sergiu Oprea
Albert García García

soprea @ dtic.ua.es
agarcia @ dtic.ua.es